

Proposal for CrisisEcho: LLM-Augmented Early Detection of Local Emergencies from Social Media

Samuel Enam Zih (Leader)

CS, Virginia Tech

samuelez@vt.edu

Christopher Drake Williams

CS, Virginia Tech

cdw24@vt.edu

Arjun Sivakumar

CS, Virginia Tech

arjuns8@vt.edu

Ki Hong Park

CS, Virginia Tech

stevep00@vt.edu

Shiyi Liu

CS, Virginia Tech

lshiyi25@vt.edu

Abstract—CrisisEcho is a consumer-facing platform that applies Retrieval-Augmented Generation (RAG) and multi-step LLM agents to detect, classify, and summarize hyperlocal emergencies from social media in near real-time. Official emergency channels routinely lag events by minutes or hours; CrisisEcho closes this gap by ingesting candidate social and official feeds, storing data across three MongoDB Atlas databases (document, geospatial, vector) plus Redis and AWS S3, and running a three-step LangChain agent chain that clusters posts, scores severity, and generates plain-language alerts. Alerts are delivered via a Next.js frontend with a Leaflet.js interactive map and WebSocket notifications, backed by a Go Fiber API. Exact data sources, backend service count, and non-core libraries are not fixed and will be finalized based on API accessibility and deployment resources available during development.

Index Terms—emergency detection, social media mining, RAG, LLM agents, MongoDB Atlas Vector Search, real-time systems, crisis informatics

I. INTRODUCTION

Official channels — 911 dispatch, government sensors, police scanners — frequently lag fast-moving crises by minutes or hours. Ordinary people post hyperlocal signals on social media well before any official report: “Water rising fast on my street!” or “Heard gunshots near the mall.” Existing tools are keyword-based and fail on indirect language, sarcasm, and regional idioms. CrisisEcho replaces rigid keyword matching with semantic RAG-powered reasoning, demonstrating four LLM-driven paradigm shifts: text to semantics, retrieval to reasoning, vertical to multi-domain, and closed-world to open-world generalization.

II. RELATED WORK

Imran et al. [1] and Olteanu et al. [2] established Twitter as a crisis sensor and built foundational lexicons, while Middleton et al. [4] showed geotagged posts correlate with ground-truth disaster events. Lewis et al. [5] introduced RAG to ground generative models in retrieved evidence, reducing hallucination. Yao et al. [6] demonstrated chained LLM reasoning steps outperform single-prompt approaches, and Karpukhin et al. [7] showed dense vector retrieval outperforms BM25 — motivating our Atlas Vector Search hybrid retrieval design.

III. DATA SOURCES

Note: The sources below are candidates. The final integrated set will be determined during development based on API rate limits, access restrictions, and available deployment resources.

Social (candidates): Twitter/X, Reddit, Bluesky, Mastodon, Nextdoor, Telegram, Facebook Groups. *Official (candidates):* NWS, USGS Earthquake, FEMA, PulsePoint, Waze Connected Citizens. *News/Crowdsourced (candidates):* GDELT, Patch.com RSS, NewsAPI, Ushahidi, Citizen App. *Training/Evaluation:* HumAID [3] (77k tweets), CrisisNLP [1], CrisisLex [2], Kaggle Disaster Tweets, CrisisMMD, FireFlood Dataset.

IV. SYSTEM ARCHITECTURE

CrisisEcho runs as a five-layer pipeline (Fig. 1). **Layer 1** platform-specific Python workers publish raw JSON to Aiven for Apache Kafka across four topics (`social_raw`, `official_alerts`, `news_feed`, `citizen_reports`). **Layer 2** a Python preprocessor applies spaCy cleaning, a cascading geocoder (Carmen → spaCy NER → location DB priors), MinHash LSH deduplication, a DistilBERT relevance filter (removes ≈70% of posts), Voyage AI text embedding (1024-dim), CLIP image embedding (512-dim), and S3 image upload. **Layer 3** stores data across five backends (Section V). **Layer 4** hybrid retrieval feeds a three-step LangChain LLM agent. **Layer 5** a Go Fiber API and Next.js frontend deliver alerts.

V. DATABASE ARCHITECTURE

MongoDB Atlas M0 (free tier) serves as the unified data platform because a single Atlas cluster natively supports document storage, 2dsphere geospatial indexing, and Atlas Vector Search — no separate vector or geospatial service required. The free tier is generous for a course project, and Atlas Vector Search supports multi-modal vector dimensions per collection.

DB 1 — Main (documents + geospatial).
Collections: `posts` (unified), `per-source_raw`
collections (`twitter_posts`, `reddit_posts`, etc.),




Fig. 1: CrisisEcho System Architecture

clusters, crises, alerts, subscriptions, official_alerts. All location fields carry a 2dsphere index. Key posts fields: post_id, source, raw_text, cleaned_text, timestamp, location (GeoJSON Point), location_confidence, text_embedding_id, image_urls (S3 links), cluster_id, crisis_type, severity_score, metadata.username_hash (SHA-256).

DB 2 — Location (geocoding support). Collections: geo_priors, place_index, location_cache (TTL-indexed, keyed by SHA-256 text hash). Enables the preprocessing cascade to avoid re-geocoding identical text.

DB 3 — Vector (Atlas Vector Search). Collections: text_embeddings (1024-dim cosine) and image_embeddings (512-dim cosine). Each document carries post_id, timestamp, lat, lng, crisis_type, is_relevant as filterable metadata inside \$vectorSearch.

DB 4 — AWS S3 (media). Raw images are uploaded to S3 (bucket/<source>/<post_id>/); only S3 URLs are stored in MongoDB, keeping documents lean.

DB 5 — Redis 7 (cache + Pub/Sub). 30-minute sliding post window, 5-minute LLM summary TTL, and alerts:live Pub/Sub channel relaying alerts to the Go WebSocket server.

VI. RAG PIPELINE AND EVALUATION

A. Hybrid Retrieval

Each 60-second trigger issues three concurrent queries: (1) Atlas Vector Search on text_embeddings (top 50 by cosine similarity, filtered by time window and bounding box); (2) MongoDB \$near on posts (geographic proximity regardless of semantics); (3) official_alerts lookup injected as high-authority context. Unresolved posts are enriched via location DB geo priors. Results are merged, deduplicated, and ranked: $0.5 \times \text{vec} + 0.3 \times \text{recency} + 0.2 \times \text{authority}$.

B. Three-Step LLM Agent

Claude claude-haiku-4-5 (primary) or Ollama Llama3 (fallback) executes: **Step 1** cluster posts into distinct events (JSON output); **Step 2** score severity 1–5 with rationale; **Step 3** generate a 2–3 sentence public alert with no PII. Publish threshold: severity ≥ 3 , ≥ 3 independent contributors from ≥ 2 platforms, geocoding confidence above minimum.

C. Evaluation

Compared against a CrisisLex keyword baseline on the same corpus. Targets: Precision@K > 0.75, Recall > 0.70, end-to-end latency < 90 s, FPR < 0.20, geocoding accuracy > 0.80, F1 improvement > 15 pp. Ground truth: HumAID/CrisisNLP test splits (offline) + PulsePoint/NWS active alerts (online).

VII. TECHNOLOGY STACK

Table I lists the confirmed core stack. Non-core libraries and auxiliary tooling are subject to change during development. Backend service count may be adjusted based on implementation decisions.

TABLE I: CrisisEcho Core Technology Stack

Component	Technology	Rationale
Backend	Go (Fiber)	High throughput; goroutine Web-Socket delivery
AI Pipeline	Python + LangChain	LLM/NLP ecosystem; LCEL agent chains
Frontend	Next.js + Leaflet.js	SSR support; geospatial map
Main DB	MongoDB Atlas M0	2dsphere geospatial; free cloud tier
Location DB	MongoDB Atlas M0	Geocoding cache + geo priors
Vector DB	Atlas Vector Search	Multi-modal ANN; no extra service
Media	AWS S3	Image storage; URLs only in MongoDB
Cache Queue	Redis 7 + Aiven Kafka	Post window; LLM TTL; WS relay Managed Kafka; decouples ingestion
LLM	Claude Haiku / Llama3	Cost-efficient; free offline fallback
Deploy	GCP Cloud Run	Serverless; free-tier prototype

REFERENCES

- [1] M. Imran, P. Mitra, C. Castillo, “Twitter as a lifeline,” *LREC*, 2016.
- [2] A. Olteanu et al., “CrisisLex,” *ICWSM*, 2014.
- [3] F. Alam et al., “HumAID,” *ICWSM*, 2021.
- [4] S. Middleton et al., “Real-time crisis mapping,” *IEEE Intell. Syst.*, vol. 29, no. 2, 2014.
- [5] P. Lewis et al., “RAG for knowledge-intensive NLP,” *NeurIPS*, 2020.
- [6] S. Yao et al., “ReAct,” *ICLR*, 2023.
- [7] V. Karpukhin et al., “Dense passage retrieval,” *EMNLP*, 2020.

TABLE II: Team Role Summary

Member	Responsibilities
S. Zih (Leader)	Backend: AI/DB/API/Docker; Frontend: map; Report: contribution
A. Sivakumar	Backend: ingestion/Kafka; Frontend: map UI — markers, pin drawer; Report: contribution
Ki Hong Park	Backend: WebSocket, NLP, Redis; Frontend: query bar, notifications; Report: contribution
Shiyi Liu	Backend: evaluation, ablations; Frontend: results display, data QC; Report: proof and submission
C. Williams	Frontend: UI/UX, dashboard, diagrams; Backend: deployment, ethics; Report: lead and integration

TABLE III: Project Schedule

Weeks	Milestone
1–2	Literature review; provision Atlas, Aiven Kafka, Redis; define schemas
3–4	Ingestion workers; spaCy + geocoding; DistilBERT fine-tune; embeddings to Atlas Vector Search
5–6	Hybrid retrieval; LangChain 3-step agent; Haiku + Llama3 fallback
7–8	Cluster persistence; Redis Pub/Sub; Go Fiber REST + WebSocket API
9–10	Next.js frontend; evaluation suite; keyword baseline comparison
11–12	Stretch features; system testing; latency and threshold tuning
13–15	Prompt refinement; documentation; demo; final report and submission

APPENDIX

All architectural and design decisions were made by the project team.